




functional view:





The quote you are asking about describes the mechanism of a **DRAM burst**, which is how modern memory hardware compensates for the incredibly slow physical reading process we discussed earlier.

Because it takes tens of nanoseconds to physically charge decoder lines, open a cell's gate, and detect its tiny electrical charge, reading data one bit at a time would severely bottleneck a modern processor 1 To solve this, DRAM relies heavily on parallel hardware design.

Here is how that process works in detail:

- **Parallel Sensing:** When a processor requests data at a specific memory location, the DRAM chip actually activates a massive row of memory cells all at once 4 5 . At the bottom of the vertical bitlines connected to these cells are thousands of sensors (sense amplifiers) 6 7 . **These sensors all work in parallel to simultaneously detect the 1s and 0s of a large chunk of consecutive memory locations** 3 .
- **The Burst Buffer:** Once the sensors have read this data, a sequence of these consecutive values is loaded into a fast, temporary storage area on the chip known as a burst buffer 8 9 .
- **High-Speed Delivery:** From this burst buffer, the data is fired off to the processor at incredible speeds 3 9 . For example, the sources describe a hardware design where 128 memory values are temporarily loaded into the burst buffer at once, and then rapidly cycled out to the processor in 16 rapid-fire sets of 8 bits 9 . **This entire sequence of rapidly delivered, consecutive data is referred to as a DRAM burst** 3 9 .

Why this matters for software performance:

This specific hardware behavior is exactly why **memory coalescing** is the single most important optimization for GPU programmers 10 11 .

If a program is written so that a warp of 32 threads requests consecutive memory locations at the exact same time, the GPU hardware combines those requests into a single, consolidated memory access 11 12 . This combined access perfectly aligns with how the DRAM chip operates: the chip only has to suffer the slow process of activating the cells and sensors once, and then it can seamlessly deliver all the data the threads need using a single, high-speed DRAM burst 3 11 .

Conversely, if threads request scattered, non-consecutive data, the program is essentially

ignoring this hardware design ¹³ ¹⁴ . The DRAM fetches a burst of consecutive data, but the threads only use a tiny fraction of it. The hardware is then forced to perform the slow, tens-of-nanoseconds sensing process over and over again for different locations, creating a massive traffic jam of memory requests that leaves the processor's compute cores sitting idle ¹⁵ ¹⁶ .

When a request for read comes in, then instead of reading 8 bits, 128 bits in sets of 8 bits are loaded into the the read burst buffer? Then, if requests are consecutive the later requests are read from the burst buffer???

Yes, your understanding is exactly correct!

Here is a detailed breakdown of how this process works based on the sources:

1. Loading the Burst Buffer When a read command is sent to a specific memory address, the hardware does not just retrieve the 8 bits you asked for. Instead, it uses part of the column address to connect 128 memory cells (and their bitlines) to the burst buffer simultaneously ¹ ² . This temporarily loads, or "caches," 128 values into the burst buffer all at once ² .

2. Rapid Delivery Once those 128 bits are sitting in the fast burst buffer, the hardware uses the remaining bits of the address to rapidly cycle through them. It sends the data to the processor in 16 quick, sequential sets of 8 bits (16 sets x 8 bits = 128 bits) ² . This rapid sequence of deliveries is known as a **burst length** of 16 ² .

3. Consecutive Requests Because the data is already queued up in the buffer, if your program's subsequent memory requests are for the consecutive data locations, they are rapidly fulfilled using this burst ² ³ . As the sources note, this allows a massive amount of data (up to 1024 bits when working across 8 chips simultaneously) to be read "extremely quickly, as long as the data is all next to one another" ³ .

This hardware design is the exact physical reason why **memory coalescing** is so crucial for software performance. When a program is written so that multiple threads request consecutive memory locations at the exact same time, the hardware combines those requests into a single consolidated request ⁴ . This allows the processor to seamlessly scoop up the consecutive data that the DRAM chip is already spitting out in its high-speed burst ⁴ ⁵ .

In tiling, where both matrices are in row-major, are the accesses naturally coalesced??

Yes, they are naturally coalesced!

When both matrices are stored in a row-major layout, the standard tiling algorithm maps perfectly to the physical memory hardware ¹ ² .


When your 2D thread block loads a tile from global memory into shared memory, consecutive threads in a warp are assigned to fetch consecutive elements across a single row of the input tile ³ . Because a row-major layout physically places elements of the same row right next to each other in DRAM, those 32 threads end up requesting a completely contiguous block of memory addresses at the exact same time ⁴ ⁵ . The GPU hardware naturally recognizes this pattern and coalesces the requests into a single, efficient DRAM burst ⁶ ⁷ .

This is exactly why the "corner turning" trick is only necessary when Matrix B is strictly stored in a column-major layout ⁸ ⁹ . If B is already row-major, your default tiled loading pattern will already be coalesced.

Since we have wrapped up the mechanics of memory coalescing, we can move on to the next topic in our plan. Would you like to explore **how the GPU hides memory**

latency using parallel channels and banks, or would you prefer to look at the Thread

Coarsening software optimization?



- This is bare metal representation of a figure 6.9.
- Our burst size is 1 element (4 bytes)
- To read element "1.2", the steps would involve long latency for the decoder to enable the cells and for the cells to share their stored charge with the sensing amplifier. During this time, the data lines are idle (shown as gray section).

When this is done and data is being sent. During this time, the data lines are being used (shown as black section)